


RoboSkate

Midterm Project Presentation

Michelle Bettendorf
Gintautas Palinauskas
Meriç Sakarya
Batuhan Yumurtacı
Finn Süberkrüb

Cloud-Based Machine Learning in Robotics
Summer Semester 2021



Agenda

- Recap of Objective
- Work done so far:
 - Current System Architecture
 - Reward Functions
 - Image Processing
 - Callback Functions
 - First Results
 - Open Issues
- Future Work

Recap of our goals



Goals:

- Setting up the simulation
- Selecting the best suited algorithm
- Teaching the agent basic movement in DummyBall
- Teaching the agent basic movement in RoboSkate
- Reaching checkpoints in game
- Finishing the whole level

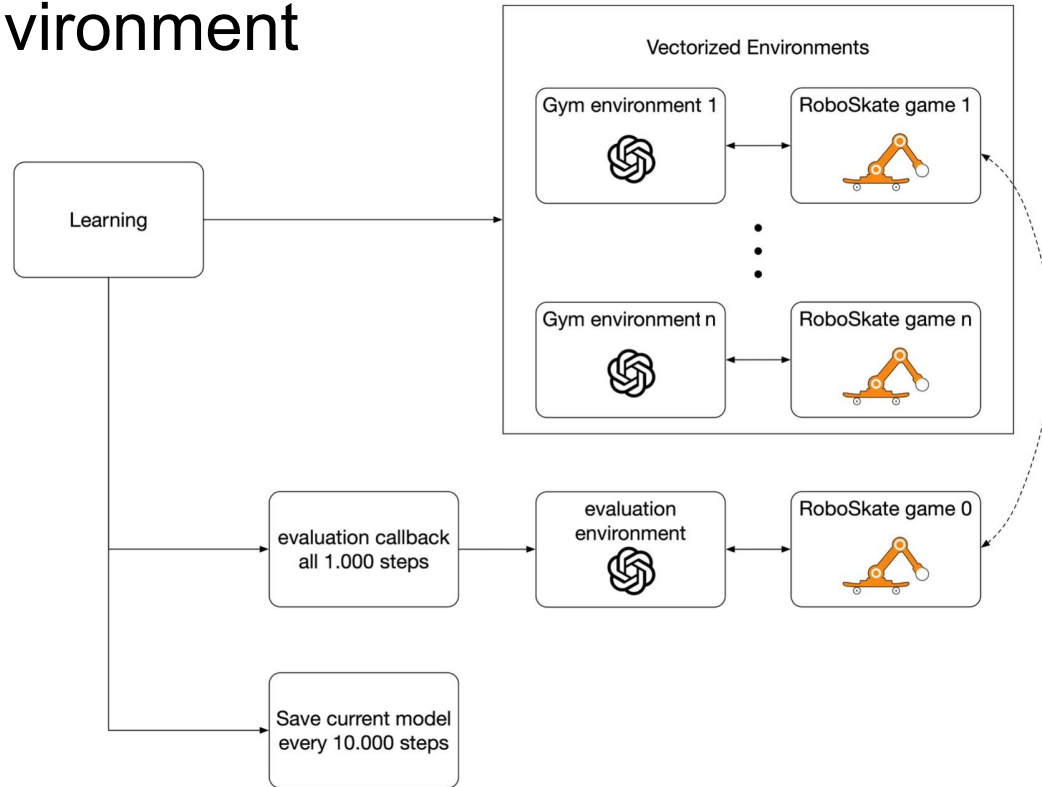
Agent




cam




Vectorized Environment




Termination




Roll angle too high



Pitch angle too high




wrong direction





Joint velocity 1 too high

Reward function



$$\text{BallHighReward} = - (\text{BallHigh} - 0.2)^4$$



First results


- 10 environments
- 20M steps
- A2C Algorithm
- MLP Policy
- 18 Observation inputs
- Reward for:
 - travelled distance
 - low ball

<https://youtu.be/PyfHF5J3TUI>


First results

- 10 environments
- 10M steps
- 28h
- A2C Algorithm
- MLP Policy
- 9 Observation inputs
- Reward for:
 - travelled distance
 - low ball
 - ball behind

<https://youtu.be/zSduDkrXD2k>




2D State Space



2D State Space

Cartesian coordinates

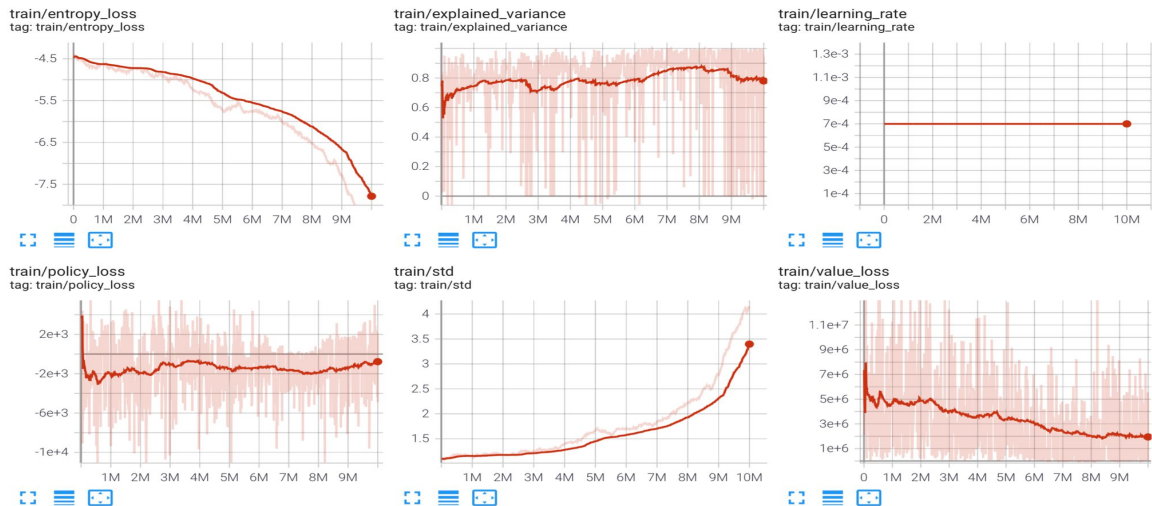


$$\begin{aligned} \text{Ball height} &= \cos(J1 + J2) L2 + \cos(J1) L1 \\ \text{Ball distance} &= \sin(J1 + J2) L2 + \sin(J1) L1 \end{aligned}$$

Tensorboard Callback

- Uses its own thread.
- Example:

```
self.logger.record('reward/mean', np.mean(self.locals.get("rewards")))
```
- Allows great visualization of the training results:




Evaluation Callback

- Evaluates the performance of the agent
- Uses separate evaluation environment
- Goal: Frequently checking up on the training to stop early
- Saves best model for further training or testing

- ```
eval_callback = EvalCallback(eval_env,
best_model_save_path="./scripts/python/RoboSkate/agent models/",
log_path="./scripts/python/RoboSkate/agent models/",
eval_freq=best agent eval_freq,
n_eval_episodes=best agent n_eval_episodes,
deterministic=True, render=False)
```

# Image Preprocessing RoboSkate

First try: Watershed




# Image Preprocessing RoboSkate


Original




Noisy




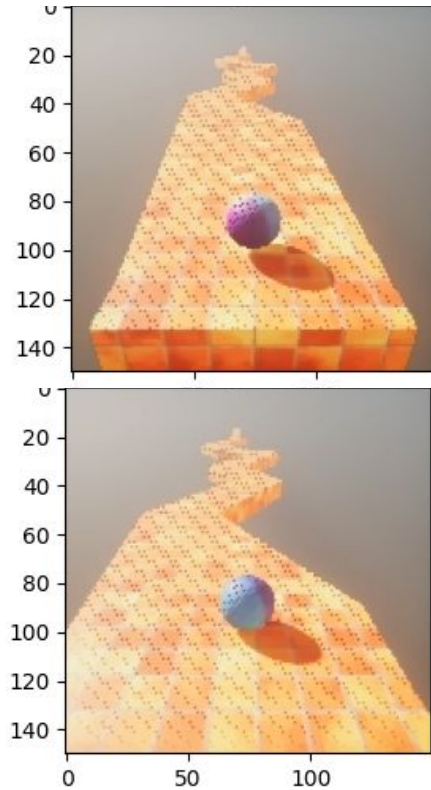
Canny filter,  $\sigma = 1$



Canny filter,  $\sigma = 3$



# Image Preprocessing: Testing on DummyBall






# Image Preprocessing: Testing on DummyBall


A2C algorithm:

location/1



Reward

Total



Canny filter



Noisy image

# DummyBall learning via images

Conditioning the observations and actions:


- Calculate z distance from platform
- Scaling action with 50

Early termination:

- When z distance < -1
- When x distance > 60

Image preprocessing using cv2:

1. Read 150x150 image
2. Convert to grayscale
3. Use Canny edge detector



# Image preprocessing ball game


Threshold method



Original image



Canny method



# Neural network architectures

# Default network architecture, from stable-baselines

if net\_arch is None:

if features\_extractor\_class == NatureCNN:

net\_arch = []

MlpPolicy

else:

net\_arch = [dict(pi=[64, 64], vf=[64, 64])]

Policy network

Value function network

150x150 image after  
convolutional layers:

1.  $(150 - 8) / 4 + 1 = 36$
2.  $(36 - 4) / 2 + 1 = 17$
3.  $(17 - 3) / 1 + 1 = 14$

```
def __init__(self, observation_space: gym.spaces.Box, features_dim: int = 512):
 super(NatureCNN, self).__init__(observation_space, features_dim)
```

```
self.cnn = nn.Sequential(
 nn.Conv2d(n_input_channels, 32, kernel_size=8, stride=4, padding=0),
 nn.ReLU(),
 nn.Conv2d(32, 64, kernel_size=4, stride=2, padding=0),
 nn.ReLU(),
 nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=0),
 nn.ReLU(),
 nn.Flatten(),
)
```

CnnPolicy

# Compute shape by doing one forward pass

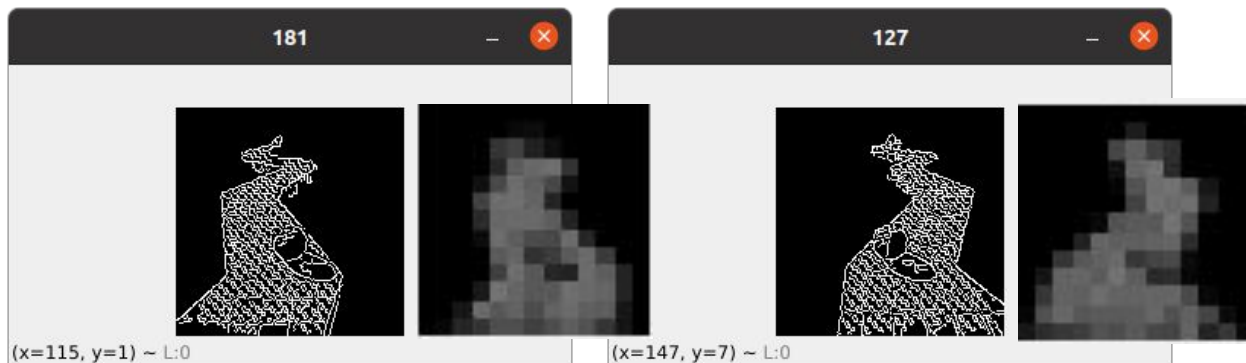
with th.no\_grad():

n\_flatten = self.cnn(th.as\_tensor(observation\_space.sample()[None]).float()).shape[1]

self.linear = nn.Sequential(nn.Linear(n\_flatten, features\_dim), nn.ReLU())


def forward(self, observations: th.Tensor) -> th.Tensor:

return self.linear(self.cnn(observations))




# Agent learning through tensorboard (PPO)

location/1  
tag: Joint/location/1



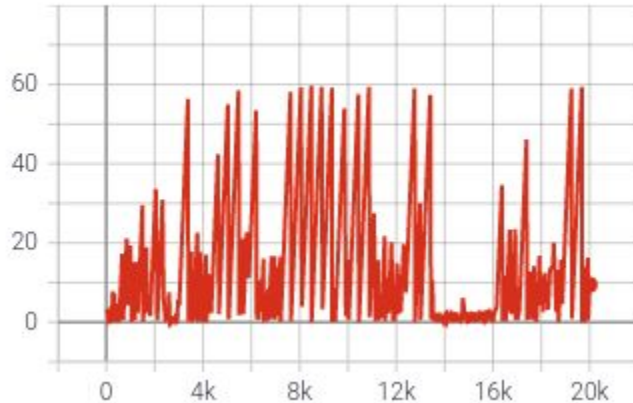
location/1  
tag: Joint/location/1



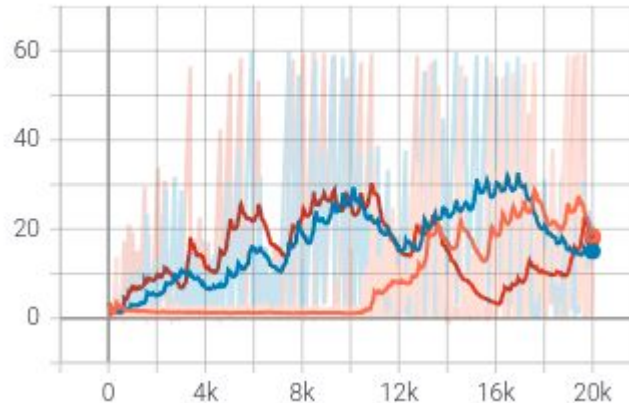
Using only default  
settings

# Agent learning through tensorboard (A2C)

location/1  
tag: Joint/location/1



location/1  
tag: Joint/location/1



Using default  
settings except  
 $n\_steps = 10$

# PPO

- Advantages:
  - Easy code
  - Sample efficient
  - Easy to tune
- On-Policy → No replay buffer
- Clipped Objective Function → Policy does not stray away

---

**Algorithm 1** PPO, Actor-Critic Style

---

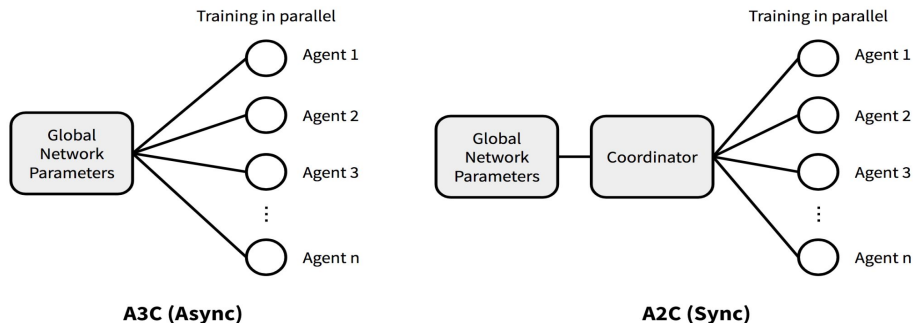
```
for iteration=1,2,... do
 for actor=1,2,...,N do
 Run policy $\pi_{\theta_{\text{old}}}$ in environment for T timesteps
 Compute advantage estimates $\hat{A}_1, \dots, \hat{A}_T$
 end for
 Optimize surrogate L wrt θ , with K epochs and minibatch size $M \leq NT$
 $\theta_{\text{old}} \leftarrow \theta$
end for
```

---

PPO Algorithm [1]

# A2C

- A3C
  - Asynchronous updates
  - Architectures share layers between the policy and value function
  - Policy gradient method
  - Multiple actors
  - Thread-specific agents → Policies of different versions → Update sub-optimal
- A2C: Synchronous version of Advantage Actor Critic (A3C)



Comparison between A2C and A3C systems




# A2C vs PPO

- PPO achieved better results
- PPO converged faster
- A2C is simpler to use
- PPO is more robust to hyper-parameter changes

PPO

location/1  
tag: Joint/location/1



A2C

location/1  
tag: Joint/location/1




# Current Issues → Possible Solutions

- Reward function too specific → Imitation Learning / Simpler environment
- NRP → Ready for DummyBall / Needs work with RoboSkate
- GPU-intensive → Simpler Environment
- Game-Physics wacky → Needs to be investigated or validated

# Simple RoboSkate Environment

Create a 2D simple Roboskate:

- Faster training (lighter)
- Faster hyperparameter optimization
- Faster reward function testing
- Enough for teaching basic movement
  - Going forward/backward
  - Accelerating
  - Brake




# Reward Function

- Needs tailoring
- Indeterminate functions are not performing well
- Specific instructions needed
- Sparsity is a huge problem → Imitation learning

# Imitation Learning

Feature and reward engineering is not always straightforward.



Hard-coded agent

Human player


Learn to perform a task from demonstrations.



# Imitation Learning



Hard-coded agent




Teleoperation demonstration  
(human player)

# Generative Adversarial Imitation Learning (GAIL)

- Model-free imitation learning algorithm
- Efficient in terms of expert data
- Uses Trust Region Policy Optimization (TRPO)
- Not sample efficient in terms of environment interactions


# Adversarial Inverse Reinforcement Learning (AIRL)

- Simultaneous learning of the reward function and value function
  - Make use of the efficient adversarial formulation
  - Recover a generalizable reward function
- Robust to changes in the environment dynamics
- Compared to GAIL
  - Similar performance in traditional imitation learning setups
  - Outperforms in transfer learning setups

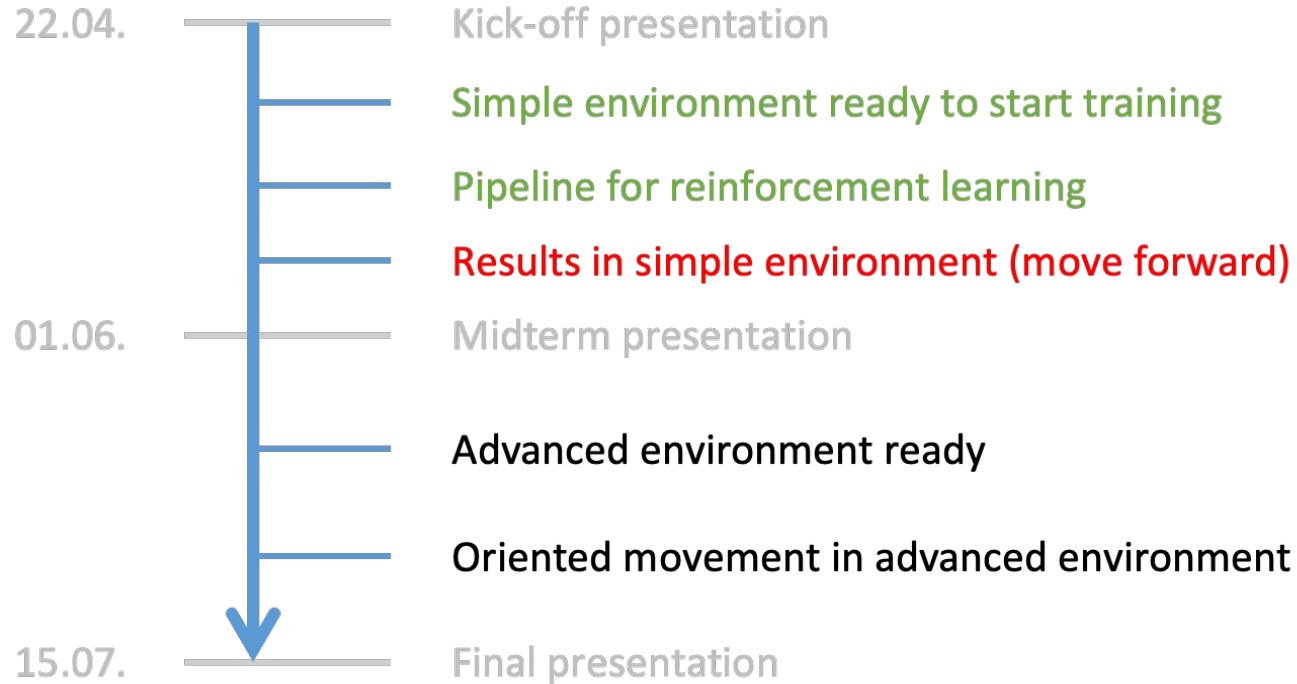




# Future work



# Schedule



# Responsibilities

|                       | Preliminary Tasks                                                                             |
|-----------------------|-----------------------------------------------------------------------------------------------|
| Michelle Bettendorf   | DummyBall, preprocessing images                                                               |
| Gintautas Palinauskas | DummyBall, preprocessing images, teleoperation, DockerFile                                    |
| Meriç Sakarya         | RoboSkate Reward functions, server training                                                   |
| Batuhan Yumurtacı     | RoboSkate Reward functions, server training                                                   |
| Finn Süberkrüb        | Vectorized Environment, Roboskate learning algorithm, Cloud administration and infrastructure |